

Fuzzing a Software Verifier

Introduction

[ESBMC](#) is an SMT-based code analyzer for single (multiple) thread C/C++ programs. It can leverage multiple solvers, including Z3, CVC4, boolector, etc. ESBMC can be used to find vulnerabilities or perform checks, like memory leak check, deadlock check, and other checks.

Fuzzing is a new software testing method to randomly generate a batch of input trying to crash the software. It is an automatic analysis method to quickly find vulnerabilities in software. Among multiple implements, [libFuzzer](#) is a state-of-the-art fuzzing module and is a part of the LLVM project.

In this task, I implemented a fuzzing framework on top of ESBMC by introducing libFuzzer. I first design new options in the ESBMC program to enable fuzzing and then compile a corresponding fuzzing program by calling libFuzzer. Then, we run the fuzzing test in one step to find vulnerabilities.

File Structures

In this directory, `Readme.md` and `Readme.pdf` are the reports of this task. The test codes and the corresponding results are located in the `test` directory. `Script` directory contains codes to compile `esbmc` and run the fuzzing test. Finally, `ESBMC_Project` contains the modified C++ codes (without Clang11), and `bin` contains the compiled binary file in both Darwin and Linux platforms.

Approaches

My experiment is implemented on a Macbook Pro. Then, I re-compile the modified ESBMC in an Ubuntu virtual machine. Finally, I integrate ESBMC's CI/CD (GitHub Actions) to confirm that our codes can be run on Windows, Linux, and Darwin environments.

Install ESBMC from source code

First, I need to install the ESBMC platform from the source code. I followed the instructions of [ESBMC](#). I used the following commands to setup dependents, and download ESBMC:

```
1 brew install gmp cmake boost ninja python3 automake && pip3 install PySMT
2 mkdir ESBMC_Project && cd ESBMC_Project && git clone https://github.com/esbmc/esbmc
```

Then, I download the clang 11 as a dependent. Note that clang provided by MacOS natively does not have libFuzzer included.

```
1 tar xJf clang+llvm-11.0.0-x86_64-apple-darwin.tar.xz && mv clang+llvm-11.0.0-x86_64-apple-darwin clang11
```

Since we do not need to modify the SMT part of ESBMC and also do not need to deal with solidity, I do not compile any solvers at first. I can start compiling ESBMC using the following commands:

```

1 cmake .. -GNinja -DBUILD_TESTING=On -DENABLE_REGRESSION=On -DBUILD_STATIC=On -
  DClang_DIR=$PWD/../../clang11 -DLLVM_DIR=$PWD/../../clang11 -
  DC2GOTO_SYSROOT=/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk -
  DCMAKE_INSTALL_PREFIX:PATH=$PWD/../../release
2
3 cmake --build . && ninja install

```

Now, the unmodified version of ESBMC should be compiled successfully.

Integrating libFuzzer

To enable ESBMC with fuzzing test, I need to enlarge options of ESBMC, I increase the following options:

```

1 $ ./ESBMC_Project/ESBMC_Project/release/bin/esbmc --help
2
3 * * *           ESBMC 6.9.0           * * *
4
5 Main Usage:
6   --input-file file.c ...           source file names
7
8 Options:
9   -? [ --help ]                   show help
10
11 Fuzzing:
12   --fuzz                           use libFuzzing to fuzzing
13   --fuzz-clang arg                 binary path to clang
14   --fuzz-sanitize arg              a list of sanitizes
15   --fuzz-coverage                  use sanitize coverage
16   --fuzz-output arg               output of fuzzing file
17   --fuzz-run arg                   runtime args
18   --fuzz-compile arg               compile args
19
20 ---

```

- `--fuzz` , enable fuzzing test
- `--fuzz-clang` , indicate the clang binary, for example: `/usr/lib/clang`
- `fuzz-sanitize` , indicate the sanitizes, for example: `address, memory, or signed-integer-overflow`
- `fuzz-coverage` , indicate whether sanitize coverage (`-fsanitize-coverage=trace-pc-guard`)
- `fuzz-output` , optional, the output file of the fuzzing binary
- `--fuzz-compile` , other arguments to compile the program, for example `-I/usr/include`
- `--fuzz-run` , arguments to run the fuzzing test, for example `-seed=3918206239` . See [LibFuzzer](#).

It is implemented by modify `esbmc/src/esbmc/options.cpp` . I added a new command group and the above commands in `all_cmd_options[]` :

```

1 ---
2 {"Fuzzing",
3  {"fuzz", NULL, "use libFuzzing to fuzzing"},

```

```

4      {"fuzz-clang",
5       boost::program_options::value<std::string>(),
6       "binary path to clang"},
7      {"fuzz-sanitize",
8       boost::program_options::value<std::string>(),
9       "a list of sanitizes"},
10     {"fuzz-coverage", NULL, "use sanitize coverage"},
11     {"fuzz-output",
12      boost::program_options::value<std::string>(),
13      "output of fuzzing file"},
14     {"fuzz-run",
15      boost::program_options::value<std::vector<std::string>>(),
16      "runtime args"},
17     {"fuzz-compile",
18      boost::program_options::value<std::vector<std::string>>(),
19      "compile args"}}},
20     ---

```

Then, I created two files `fuzzing.cpp` and `fuzzing.h` in `esbmc/src/utils`, and implemented a new class called `fuzzer`. Its declaration is in `fuzzing.h`:

```

1  class fuzzer
2  {
3  public:
4      const char *clang_path;
5      const char *common_args = "-g";
6
7      fuzzer(const char *clang_path);
8
9      ~fuzzer();
10
11     int do_fuzzing(const char *input_file);
12     int do_fuzzing(
13         const char *input_file,
14         const char *output_file,
15         const char *sanitize,
16         bool coverage);
17     int do_fuzzing(
18         const char *input_file,
19         const char *output_file,
20         const char *sanitize,
21         bool coverage,
22         const char *include,
23         const char *other,
24         const char *cmd_args);
25
26     int run_fuzz(std::string output_file, const char *cmd_args);
27 };

```

Here, `do_fuzzing` is used to compile the fuzzing test binary, and `run_fuzz` is used to run the fuzzing test. In the `do_fuzzing` function, I parsed the command arguments and then called `clang` to compile the fuzzing test binary. It parses the parameters as follows:

```
1  int fuzzer::do_fuzzing(  
2      const char *input_file,  
3      const char *output_file,  
4      const char *sanitize,  
5      bool coverage,  
6      const char *include,  
7      const char *other,  
8      const char *cmd_args)9  
10 {  
11     std::list<std::string> args = std::list<std::string>();  
12     std::string of = std::string("./a.out");  
13     if(this->clang_path == nullptr)  
14     {  
15         printf("can not find a clang binary.\n");  
16         return FUZZER_FAIL;  
17     }  
18     if(input_file == nullptr)  
19     {  
20         printf("must have a input file.\n");  
21         return FUZZER_FAIL;  
22     }  
23  
24     args.push_back(this->clang_path);  
25     args.push_back(this->common_args);  
26  
27     std::string sanitize_arg = std::string("-fsanitize=fuzzer");  
28  
29     if(sanitize != nullptr)  
30     {  
31         sanitize_arg = sanitize_arg + "," + std::string(sanitize);  
32     }  
33     args.push_back(sanitize_arg);  
34  
35     if(coverage)  
36     {  
37         std::string coverage_arg =  
38             std::string("-fsanitize-coverage=trace-pc-guard");  
39         args.push_back(coverage_arg);  
40     }  
41  
42     if(output_file != nullptr)  
43     {  
44         of = std::string(output_file);  
45         std::string output_arg = std::string("-o ") + of;
```

```

46     args.push_back(output_arg);
47 }
48
49 if(include != nullptr)
50 {
51     std::string include_arg = std::string("-I ") + std::string(include);
52     args.push_back(include_arg);
53 }
54
55 if(other != nullptr)
56 {
57     args.push_back(std::string(other));
58 }
59
60 args.push_back(input_file);
61
62 std::string cmd;
63 for(std::list<std::string>::iterator elem = args.begin(); elem != args.end();
64     elem++)
65 {
66     cmd = cmd + *elem + " ";
67 }
68 std::cout << cmd << std::endl;
69 int ret = system(cmd.c_str());
70
71 if(ret != 0)
72 {
73     return ret;
74 }
75
76 ret = fuzzer::run_fuzz(of, cmd_args);
77 return ret;
78 }

```

In the `run_fuzz` function, the compiled fuzzing test binary will be called. Its arguments are passed by `--fuze-run`. It contains the following codes:

```

1  int fuzzer::run_fuzz(std::string output_file, const char *cmd_args)
2  {
3      int ret;
4      if(output_file.size() < 1)
5      {
6          printf("wrong output file");
7          return -1;
8      }
9      if(output_file[0] != '/' && output_file[0] != '.' && output_file[0] != '~')
10     {
11         std::string run_cmd =
12             std::string("./") + output_file + " " + std::string(cmd_args);

```

```

13     std::cout << run_cmd << std::endl;
14     ret = system(run_cmd.c_str());
15 }
16 else
17 {
18     std::string run_cmd = output_file + " " + std::string(cmd_args);
19     std::cout << run_cmd << std::endl;
20     ret = system(run_cmd.c_str());
21 }
22 return ret;
23 }

```

In this way, ESBMC now provides a one-step method to run the fuzzing tests. Finally, I modified `parseoptions.cpp` and `parseoptions_baset::main()`. When the program receives a `--fuzz` flag, and then parses the options from `cmdline`:

```

1     bool coverage = false;
2     const char *output = cmdline_getval("fuzz-output");
3     const char *sanitize = cmdline_getval("fuzz-sanitize");
4     const std::list<std::string> &compile_args =
5         cmdline_get_values("fuzz-compile");
6     const std::list<std::string> &run_args = cmdline_get_values("fuzz-run");
7
8     ---
9     if(!compile_args.empty())
10    {
11        for(std::list<std::string>::const_iterator elem = compile_args.begin();
12            elem != compile_args.end();
13            elem++)
14        {
15            compile_args_full = compile_args_full + *elem;
16            compile_args_full = compile_args_full + " ";
17        }
18    }
19
20    if(!run_args.empty())
21    {
22        for(std::list<std::string>::const_iterator elem = run_args.begin();
23            elem != run_args.end();
24            elem++)
25        {
26            run_args_full = run_args_full + *elem;
27            run_args_full = run_args_full + " ";
28        }
29    }

```

After that, it will create a `fuzzer` object and call its `do_fuzzing` function to start the fuzzing test:

```




1  fuzzer f = fuzzer(cmdline_getval("fuzz-clang"));
2  ret =
3      f.do_fuzzing( input.c
4      _str(), output,
5      sanitize,
6      coverage,
7      NULL,
8      compile_args_full.c_str(),
9      run_args_full.c_str());
10 return ret;

```

Integrating CI/CD

I integrated the ESBMC's CI/CD to automatically compile modified ESBMC over platforms and check unit tests and regressions. First, I forked ESBMC's GitHub repo and activated GitHub Action. The C/C++ Lint Code based checks the C/C++ style and I modified my codes based on its suggestions.

Also, it compiles ESBMC in Linux, Darwin, and Windows platforms and run tests on those platforms. Besides, as I did not compile ESBMC with any solvers, CI/CD guarantees that all solvers work well as expected. I passed all checks in CI/CD, and the results are shown here:

33 workflow runs	Event ▾	Status ▾	Branch ▾	Actor ▾
 update c lint Lint Code Base #11: Commit 823b7ef pushed by ertuil			master	2 hours ago 1m 34s
 update c lint Build All Solvers #11: Commit 823b7ef pushed by ertuil			master	2 hours ago 1h 29m 41s
 update c lint Run clang sanitizers #11: Commit 823b7ef pushed by ertuil			master	2 hours ago 11m 32s

Fuzzing Test Result

I used the same version of the example (`fuzz_me.c`) provided by LibFuzzer. I modified it to be a pure C99 version. Its code is:

```

1  #include <stdint.h>
2  #include <stddef.h>
3
4  bool FuzzMe(const uint8_t *Data, size_t DataSize) {
5      return DataSize >= 3 &&
6          Data[0] == 'F' &&
7          Data[1] == 'U' &&
8          Data[2] == 'Z' &&
9          Data[3] == 'Z'; // :-<
10 }
11

```

```

12 extern "C" int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
13     FuzzMe(Data, Size);
14     return 0;
15 }

```

Then, I used the modified `ESBMC` to perform the fuzzing test. The only command is:

```

1 ./ESBMC_Project/ESBMC_Project/release/bin/esbmc --fuzz --fuzz-sanitize address --
fuzz-clang ./ESBMC_Project/ESBMC_Project/clang11/bin/clang++ --fuzz-compile "-L
./ESBMC_Project/ESBMC_Project/clang11/lib -L
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/lib" --fuzz-run "-
seed=3918206239" test/fuzz_me.cpp

```

Here, `--fuzz` determines to perform a fuzzing test, and I use `address` sanitizer. Here, the user should explicitly determine the path to clang. Particularly, in Macbook Pro, users need to determine the path to `libc++.so` and `libSystem.dylib`, and I use `--fuzz-compile` to complete this challenge. Finally, I use the option `--fuzz-run "-seed=3918206239"` to determine the fuzzer's arguments.

The results is:

```

1 $ ./ESBMC_Project/ESBMC_Project/release/bin/esbmc --fuzz --fuzz-sanitize address
--fuzz-clang ./ESBMC_Project/ESBMC_Project/clang11/bin/clang++ --fuzz-compile "-L
./ESBMC_Project/ESBMC_Project/clang11/lib -L
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/lib" --fuzz-run "-
seed=3918206239" test/fuzz_me.cpp
2 ./ESBMC_Project/ESBMC_Project/clang11/bin/clang++ -g -fsanitize=fuzzer,address -L
./ESBMC_Project/ESBMC_Project/clang11/lib -L
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/lib test/fuzz_me.cpp
3 a.out(92792,0x107bd0600) malloc: nano zone abandoned due to inability to
preallocate reserved vm space.
4 INFO: Seed: 3918206239
5 INFO: Loaded 1 modules (7 inline 8-bit counters): 7 [0x1007ae490, 0x1007ae497),
6 INFO: Loaded 1 PC tables (7 PCs): 7 [0x1007ae498,0x1007ae508),
7 INFO: -max_len is not provided; libFuzzer will not generate inputs larger than 4096
bytes
8 INFO: A corpus is not provided, starting from an empty corpus
9 #2 INITED cov: 3 ft: 3 corp: 1/1b exec/s: 0 rss: 38Mb
10 #3 NEW cov: 4 ft: 4 corp: 2/5b lim: 4 exec/s: 0 rss: 38Mb L: 4/4 MS: 1
CrossOver-
11 #4 REDUCE cov: 4 ft: 4 corp: 2/4b lim: 4 exec/s: 0 rss: 39Mb L: 3/3 MS: 1
EraseBytes-
12 #2425 REDUCE cov: 5 ft: 5 corp: 3/14b lim: 25 exec/s: 0 rss: 39Mb L: 10/10 MS: 1
InsertRepeatedBytes-
13 #2540 REDUCE cov: 5 ft: 5 corp: 3/11b lim: 25 exec/s: 0 rss: 39Mb L: 7/7 MS: 5
ChangeBinInt-ChangeBit-ShuffleBytes-CopyPart-EraseBytes-
14 #2651 REDUCE cov: 5 ft: 5 corp: 3/9b lim: 25 exec/s: 0 rss: 39Mb L: 5/5 MS: 1
EraseBytes-

```



```

15 #2652 REDUCE cov: 5 ft: 5 corp: 3/7b lim: 25 exec/s: 0 rss: 39Mb L: 3/3 MS: 1
EraseBytes-
16 #31325 REDUCE cov: 6 ft: 6 corp: 4/11b lim: 309 exec/s: 0 rss: 42Mb L: 4/4 MS: 3
EraseBytes-CopyPart-CMP- DE: "U\x00"-
17 #31436 REDUCE cov: 6 ft: 6 corp: 4/10b lim: 309 exec/s: 0 rss: 42Mb L: 3/3 MS: 1
EraseBytes-
18 =====
19 ==92792==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x602000163a13 at
pc 0x00010076d1da bp 0x7ff7bf795070 sp 0x7ff7bf795068
20 READ of size 1 at 0x602000163a13 thread T0
21 #0 0x10076d1d9 in FuzzMe(unsigned char const*, unsigned long) fuzz_me.cpp:9
22 #1 0x10076d23a in LLVMFuzzerTestOneInput fuzz_me.cpp:13
23 #2 0x100788d80 in fuzzer::Fuzzer::ExecuteCallback(unsigned char const*,
unsigned long) FuzzerLoop.cpp:559
24 #3 0x1007884c5 in fuzzer::Fuzzer::RunOne(unsigned char const*, unsigned long,
bool, fuzzer::InputInfo*, bool*) FuzzerLoop.cpp:471
25 #4 0x100789c11 in fuzzer::Fuzzer::MutateAndTestOne() FuzzerLoop.cpp:702
26 #5 0x10078a695 in fuzzer::Fuzzer::Loop(std::__1::vector<fuzzer::SizedFile,
fuzzer::fuzzer_allocator<fuzzer::SizedFile> >&) FuzzerLoop.cpp:838
27 #6 0x1007785d2 in fuzzer::FuzzerDriver(int*, char***, int (*)(unsigned char
const*, unsigned long)) FuzzerDriver.cpp:847
28 #7 0x1007a4b92 in main FuzzerMain.cpp:20
29 #8 0x107b5551d in start+0x1cd (dyld:x86_64+0x551d)
30
31 0x602000163a13 is located 0 bytes to the right of 3-byte region
[0x602000163a10,0x602000163a13)
32 allocated by thread T0 here:
33 #0 0x100b639dd in wrap_Znam+0x7d
(libclang_rt.asan_osx_dynamic.dylib:x86_64h+0x519dd)
34 #1 0x100788c91 in fuzzer::Fuzzer::ExecuteCallback(unsigned char const*,
unsigned long) FuzzerLoop.cpp:544
35 #2 0x1007884c5 in fuzzer::Fuzzer::RunOne(unsigned char const*, unsigned long,
bool, fuzzer::InputInfo*, bool*) FuzzerLoop.cpp:471
36 #3 0x100789c11 in fuzzer::Fuzzer::MutateAndTestOne() FuzzerLoop.cpp:702
37 #4 0x10078a695 in fuzzer::Fuzzer::Loop(std::__1::vector<fuzzer::SizedFile,
fuzzer::fuzzer_allocator<fuzzer::SizedFile> >&) FuzzerLoop.cpp:838
38 #5 0x1007785d2 in fuzzer::FuzzerDriver(int*, char***, int (*)(unsigned char
const*, unsigned long)) FuzzerDriver.cpp:847
39 #6 0x1007a4b92 in main FuzzerMain.cpp:20
40 #7 0x107b5551d in start+0x1cd (dyld:x86_64+0x551d)
41
42 SUMMARY: AddressSanitizer: heap-buffer-overflow fuzz_me.cpp:9 in FuzzMe(unsigned
char const*, unsigned long)
43 Shadow bytes around the buggy address:
44 0x1c040002c6f0: fa fa fd fa fa fa fd fa fa fa fd fa fa fa fd fa
45 0x1c040002c700: fa fa fd fa fa fa fd fa fa fa fd fa fa fa fd fa
46 0x1c040002c710: fa fa fd fa fa fa fd fa fa fa fd fd fa fa fd fa
47 0x1c040002c720: fa fa fd fa fa fa fd fa fa fa fd fa fa fa fd fa
48 0x1c040002c730: fa fa fd fa fa fa fd fa fa fa fd fd fa fa fd fa

```

```

49 =>0x1c040002c740: fa fa[03]fa fa fa fa fa fa fa fa fa fa fa fa
50 0x1c040002c750: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
51 0x1c040002c760: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
52 0x1c040002c770: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
53 0x1c040002c780: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
54 0x1c040002c790: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
55 Shadow byte legend (one shadow byte represents 8 application bytes):
56 Addressable:                00
57 Partially addressable: 01 02 03 04 05 06 07
58 Heap left redzone:          fa
59 Freed heap region:           fd
60 Stack left redzone:          f1
61 Stack mid redzone:           f2
62 Stack right redzone:         f3
63 Stack after return:          f5
64 Stack use after scope:       f8
65 Global redzone:              f9
66 Global init order:           f6
67 Poisoned by user:            f7
68 Container overflow:          fc
69 Array cookie:                ac
70 Intra object redzone:        bb
71 ASan internal:               fe
72 Left alloca redzone:         ca
73 Right alloca redzone:        cb
74 Shadow gap:                  cc
75 ==92792==ABORTING
76 MS: 1 ChangeByte-; base unit: 1c12b63a941811b8a4940d6faa75a377401162a4
77 0x46,0x55,0x5a,
78 FUZ
79 artifact_prefix='./'; Test unit written to ./crash-
80 0eb8e4ed029b774d80f2b66408203801cb982a60
81 Base64: RIVa

```

As it shows that ESBMC finds a heap buffer overflow vulnerability on this program.

Conclusion

In this task, I merged the fuzzing test function into ESBMC. I used LibFuzzer to provide the fuzzing capabilities. Also, I ran the ESBMC fuzzing test targeting an example C file and found the vulnerability as expected. Finally, I integrated ESBMC's CI/CD to guarantee the code quality, cross-platform, and multiple solvers abilities.